

# attentionUnet\_10

June 15, 2023

```
[ ]: ###  
!pip3 install albumentations
```

```
[13]: import numpy as np  
import os  
from torch.utils.data import Dataset  
import torch  
from PIL import Image  
import matplotlib.pyplot as plt  
from albumentations.pytorch import ToTensorV2  
import albumentations as A  
import torch.nn.functional as F  
import torch.nn as nn  
from torch.optim import Adam  
from tqdm import tqdm, trange
```

```
[3]: ###  
class carlaData(Dataset):  
    def __init__(self, img_dir, transform = None):  
        self.transforms = transform  
        image_paths = [i+'/'+'CameraRGB' for i in img_dir]  
        seg_paths = [i+'/'+'CameraSeg' for i in img_dir]  
        self.images, self.masks = [], []  
        for i in image_paths:  
            imgs = os.listdir(i)  
            self.images.extend([i+'/'+'img' for img in imgs])  
        for i in seg_paths:  
            masks = os.listdir(i)  
            self.masks.extend([i+'/'+'mask' for mask in masks])  
    def __len__(self):  
        return len(self.images)  
    def __getitem__(self, index):  
        img = np.array(Image.open(self.images[index]))  
        mask = np.array(Image.open(self.masks[index]))  
        if self.transforms is not None:  
            aug = self.transforms(image=img, mask=mask)  
            img = aug['image']
```

```

        mask = aug['mask']
        mask = torch.max(mask,dim=2)[0]
        return img,mask
    """

```

```

[27]: data_dir = ['archive' + '/data'+i+'/data'+i for i in ['A','B','C','D','E']]
    """
    t1 = A.Compose([
        A.Resize(128,128),
        A.augmentations.transforms.Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.
    ↪5)),
        ToTensorV2()
    ])
    """

    def get_images(image_dir,transform = ↵
    ↪None,batch_size=1,shuffle=True,pin_memory=True):
        data = carlaData(image_dir,transform = t1)
        train_size = int(0.8 * data.__len__())
        test_size = data.__len__() - train_size
        train_dataset, test_dataset = torch.utils.data.random_split(data,↵
    ↪[train_size, test_size])
        train_batch = torch.utils.data.DataLoader(train_dataset,↵
    ↪batch_size=batch_size, shuffle=shuffle, pin_memory=pin_memory)
        test_batch = torch.utils.data.DataLoader(test_dataset,↵
    ↪batch_size=batch_size, shuffle=shuffle, pin_memory=pin_memory)
        return train_batch,test_batch
    """

    train_batch,test_batch = get_images(data_dir,transform =t1,batch_size=1)
    """

    class SelfAttention(nn.Module):
        def __init__(self, in_dim):
            super(SelfAttention, self).__init__()
            self.chanel_in = in_dim

            self.query_conv = nn.Conv2d(in_channels = in_dim, out_channels = in_dim/
    ↪8, kernel_size = 1)
            self.key_conv = nn.Conv2d(in_channels = in_dim, out_channels = in_dim//
    ↪8, kernel_size = 1)
            self.value_conv = nn.Conv2d(in_channels = in_dim, out_channels = ↵
    ↪in_dim, kernel_size = 1)

            self.gamma = nn.Parameter(torch.zeros(1))

        def forward(self, x):
            m_batchsize, C, width, height = x.size()

```

```

        proj_query = self.query_conv(x).view(m_batchsize, -1, width*height).
        ↪permute(0, 2, 1)
        proj_key = self.key_conv(x).view(m_batchsize, -1, width*height)
        energy = torch.bmm(proj_query, proj_key)
        attention = torch.softmax(energy, dim = -1)
        proj_value = self.value_conv(x).view(m_batchsize, -1, width*height)

        out = torch.bmm(proj_value, attention.permute(0, 2, 1))
        out = out.view(m_batchsize, C, width, height)

        out = self.gamma*out + x
        return out, attention

```

```

[15]: class DoubleConv(nn.Module):
        """(convolution => [BN] => ReLU) * 2"""

        def __init__(self, in_channels, out_channels, mid_channels=None):
            super().__init__()
            if not mid_channels:
                mid_channels = out_channels
            self.double_conv = nn.Sequential(
                nn.Conv2d(in_channels, mid_channels, kernel_size=3, padding=1, ↪
        ↪bias=False),
                nn.BatchNorm2d(mid_channels),
                nn.ReLU(inplace=True),
                nn.Conv2d(mid_channels, out_channels, kernel_size=3, padding=1, ↪
        ↪bias=False),
                nn.BatchNorm2d(out_channels),
                nn.ReLU(inplace=True)
            )

        def forward(self, x):
            return self.double_conv(x)

class Down(nn.Module):
        """Downscaling with maxpool then double conv"""

        def __init__(self, in_channels, out_channels):
            super().__init__()
            self.maxpool_conv = nn.Sequential(
                nn.MaxPool2d(2),
                DoubleConv(in_channels, out_channels)
            )

        def forward(self, x):
            return self.maxpool_conv(x)

class Up(nn.Module):
        """Upscaling then double conv"""

```

```

def __init__(self, in_channels, out_channels, bilinear=True):
    super().__init__()

    # if bilinear, use the normal convolutions to reduce the number of
    ↪ channels
    if bilinear:
        self.up = nn.Upsample(scale_factor=2, mode='bilinear', ↪
    ↪ align_corners=True)
        self.conv = DoubleConv(in_channels, out_channels, in_channels // 2)
    else:
        self.up = nn.ConvTranspose2d(in_channels, in_channels // 2, ↪
    ↪ kernel_size=2, stride=2)
        self.conv = DoubleConv(in_channels, out_channels)

def forward(self, x1, x2):
    x1 = self.up(x1)
    # input is CHW
    diffY = x2.size()[2] - x1.size()[2]
    diffX = x2.size()[3] - x1.size()[3]

    x1 = F.pad(x1, [diffX // 2, diffX - diffX // 2,
                    diffY // 2, diffY - diffY // 2])

    x = torch.cat([x2, x1], dim=1)
    return self.conv(x)

class OutConv(nn.Module):
    def __init__(self, in_channels, out_channels):
        super(OutConv, self).__init__()
        self.conv = nn.Conv2d(in_channels, out_channels, kernel_size=1)

    def forward(self, x):
        return self.conv(x)

```

```

[21]: """
device = "cuda" if torch.cuda.is_available() else "cpu"
class unet_model(nn.Module):
    def __init__(self, n_channels=64, n_classes=23, bilinear=False):
        super(unet_model, self).__init__()
        self.n_channels = n_channels
        self.n_classes = n_classes
        self.bilinear = bilinear

        self.inc = (DoubleConv(3, 64))
        self.down1 = (Down(64, 128))

```

```

self.down2 = (Down(128, 256))
self.down3 = (Down(256, 512))
self.attention_1 = SelfAttention(64)
self.attention_2 = SelfAttention(128)
self.attention_3 = SelfAttention(256)
self.attention_4 = SelfAttention(512)

factor = 2 if bilinear else 1
self.down4 = (Down(512, 1024 // factor))
self.up1 = (Up(1024, 512 // factor, bilinear))
self.up2 = (Up(512, 256 // factor, bilinear))
self.up3 = (Up(256, 128 // factor, bilinear))
self.up4 = (Up(128, 64, bilinear))
self.outc = (OutConv(64, n_classes))

def forward(self, x):
    x1 = self.inc(x)
    x2 = self.down1(x1)
    x3 = self.down2(x2)
    x4 = self.down3(x3)
    x5 = self.down4(x4)
    v1, _ = self.attention_4(x4)
    v2, _ = self.attention_3(x3)
    v3, _ = self.attention_2(x2)
    v4, _ = self.attention_1(x1)
    x = self.up1(x5, v1)
    x = self.up2(x, v2)
    x = self.up3(x, v3)
    x = self.up4(x, v4)
    logits = self.outc(x)
    return logits

```

###

```

[ ]: DEVICE = "cuda" if torch.cuda.is_available() else "cpu"

###
model = unet_model().to(DEVICE)
LEARNING_RATE = 1e-4
num_epochs = 40
loss_fn = nn.CrossEntropyLoss()
optimizer = Adam(model.parameters(), lr=LEARNING_RATE)
scaler = torch.cuda.amp.GradScaler()
###
for epoch in range(num_epochs):
    loop = tqdm(enumerate(train_batch), total=len(train_batch))
    for batch_idx, (data, targets) in loop:
        data = data.to(DEVICE)

```

```

        targets = targets.to(DEVICE)
        targets = targets.type(torch.long)
        # forward
        with torch.cuda.amp.autocast():
            predictions = model(data)
            loss = loss_fn(predictions, targets)
        # backward
        optimizer.zero_grad()
        scaler.scale(loss).backward()
        scaler.step(optimizer)
        scaler.update()

        # update tqdm loop
        loop.set_postfix(loss=loss.item())
# %%
def check_accuracy(loader, model):
    num_correct = 0
    num_pixels = 0
    dice_score = 0
    model.eval()

    with torch.no_grad():
        for x, y in loader:
            x = x.to(DEVICE)
            y = y.to(DEVICE)
            softmax = nn.Softmax(dim=1)
            preds = torch.argmax(softmax(model(x)),axis=1)
            num_correct += (preds == y).sum()
            num_pixels += torch.numel(preds)
            dice_score += (2 * (preds * y).sum()) / ((preds + y).sum() + 1e-8)

    print(f"Got {num_correct}/{num_pixels} with acc {num_correct/num_pixels*100:
↵.2f}")
    print(f"Dice score: {dice_score/len(loader)}")
    model.train()
# %%
check_accuracy(train_batch, model)
# %%
check_accuracy(test_batch, model)

```

```

[26]: #%%
for x,y in test_batch:
    x = x.to(DEVICE)
    fig , ax = plt.subplots(3, 3, figsize=(18, 18))
    softmax = nn.Softmax(dim=1)
    preds = torch.argmax(softmax(model(x)),axis=1).to('cpu')
    img1 = np.transpose(np.array(x[0,:,:,:].to('cpu')),(1,2,0))

```

```


preds1 = np.array(preds[0,:,:])
mask1 = np.array(y[0,:,:])
img2 = np.transpose(np.array(x[0,:,:,:].to('cpu')),(1,2,0))
preds2 = np.array(preds[0,:,:])
mask2 = np.array(y[0,:,:])
img3 = np.transpose(np.array(x[0,:,:,:].to('cpu')),(1,2,0))
preds3 = np.array(preds[0,:,:])
mask3 = np.array(y[0,:,:])
ax[0,0].set_title('Image')
ax[0,1].set_title('Prediction')
ax[0,2].set_title('Mask')
ax[1,0].set_title('Image')
ax[1,1].set_title('Prediction')
ax[1,2].set_title('Mask')
ax[2,0].set_title('Image')
ax[2,1].set_title('Prediction')
ax[2,2].set_title('Mask')
ax[0][0].axis("off")
ax[1][0].axis("off")
ax[2][0].axis("off")
ax[0][1].axis("off")
ax[1][1].axis("off")
ax[2][1].axis("off")
ax[0][2].axis("off")
ax[1][2].axis("off")
ax[2][2].axis("off")
ax[0][0].imshow(img1)
ax[0][1].imshow(preds1)
ax[0][2].imshow(mask1)
ax[1][0].imshow(img2)
ax[1][1].imshow(preds2)
ax[1][2].imshow(mask2)
ax[2][0].imshow(img3)
ax[2][1].imshow(preds3)
ax[2][2].imshow(mask3)
break

```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



[ ]:

[ ]: